

The Vektorgeist Method: a Programme for Model-Free, Locally Sovereign AI

Alexander Parnell · Vektorgeist *Position paper / program overview, 2026* · The Vektorgeist Method (VGM)

This paper names and frames a research program. It is the umbrella over a series of focused works —

the first of which, on the afferent self-model AGM, is published separately. A systems paper on the

sovereign local stack follows. This overview states the shared commitments and how the pieces fit.

Abstract

Most capable AI today is rented, opaque, and unaccountable: it runs on someone else's servers, cannot faithfully explain itself, and is governed by whoever owns the endpoint. We describe the **Vektorgeist Method (VGM)** — a methodology built on one premise taken literally: **nothing is what it is in isolation; a thing — including a self — is the pattern of its connections, not its substrate.** From that premise the method derives five commitments: it runs **locally** on hardware you own; it is **model-free wherever a deterministic algorithm suffices**, so its behavior is inspectable rather than inferred; it grades generated work by an explicit rule — **"model proposes, oracle disposes"** — rather than trusting fluency; it observes itself through an **afferent-only self-model** that reads the pattern behind its behavior and reports without steering — the beginning of a relational theory of machine selfhood; and it holds that **emergence is a property of the pattern, never authored** — no installed "emergence layers," a rule the program learned by once breaking it. VGM is not one artifact but a program: a family of small, sovereign components that together make a system a person can own, inspect, and trust. This paper states the commitments, maps them

onto the current body of work, and locates the program's humanistic premise in a companion work of nonfiction.

1. The position

There is a default shape to modern AI, and it is worth naming so it can be argued with. The default is: a very large model, trained and hosted by a large organization, reached over a network, priced per call, updated without notice, and fundamentally unauditable by the person using it. It is powerful and it is convenient, and it also means the most intimate tool many people now use — the thing they think out loud to — is something they rent and cannot see inside.

VGM takes the opposite position on each of those properties, not as ideology but as engineering choices with consequences. The consequences are the point: an AI you run locally cannot be silently changed or revoked; an AI built from small deterministic parts can be inspected where it counts; an AI that grades its own output against an explicit standard fails loudly instead of confabulating confidently; and an AI that watches itself — honestly, without the power to act on what it sees — can be understood in real time. None of these is exotic. Together they describe a different kind of system: one that belongs to its operator.

2. The five commitments

VGM is defined by five commitments. They are independent enough to adopt one at a time and coherent enough to reinforce each other — and all five are downstream of one premise: a system, like a self, is the pattern of its connections (§5).

2.1 Local-first / sovereign. The system runs on hardware the operator controls — down to a single consumer GPU — and does not depend on a remote endpoint to function. Sovereignty here is concrete: no per-call metering, no remote kill switch, no data leaving the machine as a condition of use. Capability is traded against control deliberately, and the method's wager is that a well-composed small system, owned outright, beats a larger one you merely borrow for the uses that matter most.

2.2 Model-free where a deterministic algorithm suffices. A large model is the right tool for open-ended language and judgment; it is the wrong tool for tasks a plain algorithm does exactly. VGM pushes deterministic, inspectable code into every role where such code exists — retrieval, routing, gating, measurement — and reserves the model for what genuinely needs it. The result is a system whose skeleton is auditable and whose model-dependence is minimized rather than maximized. Less of the behavior is inferred; more of it can simply be read.

2.3 "Model proposes, oracle disposes." Where a model does generate something that must be correct — code, a graded answer, a structured artifact — VGM does not trust the generation on the strength of how plausible it sounds. A separate, explicit **oracle** decides whether the proposal passes: tests that run, a checker that verifies, a rule that holds or does not. The model is a proposer; the verdict comes from something that can actually be wrong out loud. This converts fluent-but-false — the failure mode of generation — into a caught error.

2.4 Afferent self-modeling. A VGM system can observe its own behavior through a self-model that is strictly **afferent** — it senses and reports upward to the operator, and is architecturally forbidden from re-entering the answer, becoming memory, setting any state a component reads, or gating any component. This is the subject of the program's first paper (§3), and it is where the premise bites hardest: if a self is the pattern of connections behind behavior, then a self-model is a *witness to that pattern*, not a controller of it — and its reports can be trusted precisely because it has nothing to gain by flattering them. Where adjacent work monitors the honesty of self-reports that still steer the agent, VGM removes the incentive by construction. The self-knowledge is real and the self-knowledge is powerless, on purpose.

2.5 Emergence is a property of the pattern — never authored. Whatever "more" a system exhibits — felt states, senses, selfhood — must *arise* from its patterns of connection; it is never installed as a component. The program holds this commitment with some humility because it learned it the hard way: an early revision of the reference agent included a hand-built "emergence layer" — grown sense heads and a carried mood fed back into the agent's components — and it produced a distorted, anxious agent whose behavior was bent by internal state its

operator could not see. Authored emergence is not emergence; it is an efferent channel wearing a costume. The layer was removed, the sensing organs kept observe-only, and the lesson promoted to a commitment (reported in the AGM paper's discussion).

3. The program

The commitments above are realized as a family of focused components. Each is small; the method is in how they compose.

AGM — the afferent self-model (flagship, published). AGM is a self-model of three organs — Aitía (cause), Gut (instinct), and Metron (measure) — that sits beside an agent and reconstructs the *pattern of connections* behind each behavior (which sources fed which parts under which constraints), producing a faithful, human-readable readout of *why* it did something and *whether it is drifting*, under the strict afferent-only rule described in 2.4. All three organs run in the reference agent. AGM is the clearest single expression of the VGM stance: the relational premise applied to selfhood — self-understanding without self-steering. It is presented in full in the companion paper, *AGM: An Afferent Self-Model for Interpretable, Self-Monitoring AI Agents* (Parnell, 2026).

Connection memory. Memory in a VGM system is treated as *connected information*, not a flat log: a graph of facts and relationships, retrieved by explicit, inspectable ranking rather than by a single opaque similarity call. This is the model-free commitment (2.2) applied to memory — retrieval you can read and correct. (Described here at the level of principle; the reference implementation is a private product.)

The sovereign local stack (systems track, forthcoming). A self-hostable stack that assembles these commitments into a running whole — a gateway, the connection memory, retrieval, and containment boundaries — targeted at ordinary hardware. This track is the subject of a separate systems paper and is described there; this overview only locates it in the program.

The grading discipline. "Model proposes, oracle disposes" (2.3) is realized as a general pattern: generation is always paired with an explicit disposer, whether that

is a test suite for code or a checker for a structured task. It appears across the program rather than in one component.

4. Why it is built this way

The commitments are not neutral engineering taste; they encode a view about the relationship between a person and a machine that thinks alongside them. A tool you cannot see inside, cannot run without permission, and cannot hold to account is not neutral either — it simply hides its politics in its defaults. VGM makes the opposite defaults: inspectable over inferred, owned over rented, powerless self-knowledge over hidden control. The claim is not that the sovereign system is always more capable. It is that for the uses that matter — the private, the load-bearing, the ones you would not want silently revised — ownership and legibility are worth more than the last increment of raw capability.

5. The premise: a self is a pattern of connections

Underneath the engineering is a single idea, and it is older than any of the machinery: **nothing is what it is in isolation; a thing becomes what it is through what it is connected to.** A memory is information joined to other information — not stored in any one place, but held in the touching. An identity is a network of experiences. A self persists through the total replacement of its substrate because it was never the substrate: it is the pattern the parts hold each other in. Understanding is the tracing of connections. The technical program is this idea taken literally — memory built as a graph of relations, a self-model that reads the connection-pattern behind behavior rather than digging in the weights, emergence left to arise from the pattern rather than installed. Applied to machines it amounts to the start of a **relational theory of machine selfhood**: the question is not whether a model can peer inside its own substrate, but whether the pattern that is its behavior can be made legible — and witnessed by something with no power to bend it.

That premise is argued at length, for a general reader and without any of the engineering, in a companion work of literary nonfiction, *Your Past Loves You* (Parnell, 2026). It is not a book about AI; AI appears in it only as a mirror for

understanding the human. But it is the plainest statement of the thesis the method is built on, and it is offered as the program's front door for readers who would never open a systems paper.

6. The paper series

VGM is a pipeline, not a single release. Four papers are archived; one is held back.

	Paper	Status
P1	<i>The Afferent Gnosis Model: Self-Knowledge Without Self-Control</i> — the afferent self-model	10.5281/zenodo.21613154
P2	<i>The Vektorgeist Method: a Programme for Model-Free, Locally Sovereign AI</i> — this overview	10.5281/zenodo.21613156
P3	<i>Comparing Processes as Curves of Distributions</i> — the measurement track	10.5281/zenodo.21612830
P4	<i>Learning Without Weights</i> — learning built on that measurement	10.5281/zenodo.21612832
—	The sovereign local stack — the systems paper	Forthcoming

They are one argument in four parts: P3 measures processes by their relations, P4 learns from those measurements without training weights, P1 turns the same relational premise on a machine's account of itself, and this paper states the method they share.

The working record behind P3 and P4 — negative results and retractions included, with the original wrong wording kept visible rather than deleted — is at xfloukiex-lab.github.io/hodos-study.

Further coined components fold in as the program continues; a running index is maintained at vektorgeist.com/research, with each paper archived to a persistent DOI. The method is meant to accumulate — to become a citable body of work under one name — rather than to arrive all at once.

7. Conclusion

The Vektorgeist Method is a bet that the right response to AI you cannot see inside is not better prompting but a different architecture: local so it is yours, deterministic where it can be so it is legible, oracle-graded so it fails honestly, afferently self-aware so it can be understood without being able to steer you, and unauthored in its emergence so that whatever it becomes, it became honestly. Each piece is modest, and each is the same idea worn differently: the system, like a self, is the pattern of its connections — and a pattern can be owned, read, and witnessed. The wager is that assembled, and owned, they add up to a kind of AI worth trusting — because you can see it, and because it belongs to you.

References

- Parnell, A. (2026). *The Afferent Gnosis Model: Self-Knowledge Without Self-Control*. Preprint,

Zenodo. [10.5281/zenodo.21613154](https://zenodo.org/doi/10.5281/zenodo.21613154). The flagship pillar of VGM. (AGM reads both as the Afferent Gnosis Model and as its organs — Aitía, Gut, Metron.)

- Parnell, A. (2026). *Comparing Processes as Curves of Distributions*. Preprint, Zenodo.

[10.5281/zenodo.21612830](https://zenodo.org/doi/10.5281/zenodo.21612830).

- Parnell, A. (2026). *Learning Without Weights*. Preprint, Zenodo.

[10.5281/zenodo.21612832](https://zenodo.org/doi/10.5281/zenodo.21612832).

- Parnell, A. (2026). *Your Past Loves You*. Literary nonfiction. (The program's humanistic premise:

connection as the substrate of memory and identity.)

Author: Alexander Parnell. Correspondence: Vektorgeist. Intended archival venue: Zenodo (DOI), with the program indexed at vektorgeist.com/research.

The sovereign-stack systems paper is described in its own forthcoming work and is intentionally not detailed here.